

Zixuan Li

lizx2122@mails.jlu.edu.cn | +86 188 3792 0249 | Hong Kong, China

EDUCATION

The Hong Kong University of Science and Technology

08/2026 – 06/2027 (Incoming)

M.Sc. in Big Data Technology

Hong Kong, China

Jilin University

09/2022 – 07/2026

B.Eng. in Computer Science and Technology

Changchun, China

• GPA: 88.79/100 (Top 13%)

• Core Coursework: Computer Architecture (96), Database & Security (A+), Algorithm Design (93.7), Compiler Principles (93.2), Artificial Intelligence, Operating Systems.

PUBLICATIONS

[1] **FLINT: Influence-Guided Active Learning Framework for LoRA via Curvature-Aware Data Selection and Fine-Tuning.** *Z. Li, S. Guo, et al.*; Accepted to the Conference on Language Modeling (COLM) 2026.

[2] **Multimodal Fake News Detection using Graph Neural Networks and Attention Mechanisms.** *Z. Li*; *Advances in Engineering Innovation*, 2024

RESEARCH EXPERIENCE

FLINT — Influence-Guided Active Learning for LoRA Fine-Tuning

12/2025 – 04/2026

Jilin University / Research Project | **First Author**

- Designed **FLINT**, a two-stage active learning pipeline integrating entropy pre-filtering, **SVD** cold-start stabilization, and **K-FAC** curvature-aware influence scoring optimized for the LoRA manifold.
- Streamlined candidate evaluations by **70–80%** via adaptive entropy thresholding, successfully constraining the data selection to a strict **20%** budget across three active learning iterations.
- Achieved LLaMA-3.2-3B accuracies of **71.13%** / **66.8%** / **73.0%** on GSM8K / BBH / StrategyQA, consistently outperforming full-data LoRA baselines (**68.85%** / **64.14%** / **69.0%**).
- Accelerated end-to-end selection and training to **1.67h** on a single V100 GPU, yielding a **5.73×** speedup over LESS and **2.80×** over DataInf with a peak VRAM of **16.6GB**.
- **Accepted to COLM 2026** as first author.

Dimension-Mapped Near-Memory Computing for Sparse Attention Acceleration

05/2025 – 12/2025

Jilin University HPC Center | **First Author (Advisor: Prof. Xiaohui Wei)**

- Proposed a novel **dimension-based dataflow mapping** (as opposed to standard token-level partitioning) for Near-Memory Processing (NMP) systems, enforcing strict load balancing across parallel DRAM bank arrays via hash-routed feature decoupling.
- Developed a trace-driven hardware-software co-simulation engine using PyTorch hooks and SimPy discrete-event modeling to simulate physical DRAM row-buffer conflicts under highly skewed, long-tail sparse attention masks.
- Achieved a **27.7%** reduction in end-to-end LLM inference step latency (6920 ns → 5000 ns) and minimized workload variance (coefficient of variation dropped from **23.89%** → **7.09%**).

Multimodal Fake News Detection via GNN and Cross Attention Mechanisms

12/2024 – 02/2025

Jilin University | **First Author**

- Formulated a multimodal fusion network architecture combining BERT (textual), CNN+CLIP (visual), and GraphSAGE (social graph structure) with an adaptive cross-modal attention routing mechanism.
- Demonstrated robust performance gains over baseline models on FakeNewsNet and Sina Weibo benchmarks.
- Independently conceptualized the methodology, performed extensive ablation studies, and drafted the full manuscript published in *Advances in Engineering Innovation*.

INDUSTRY EXPERIENCE

Tencent — WeChat Headquarters (WeChat Agent Research)

05/2026 – Present

Research Intern

- Algorithmic Optimization for Multi-Turn Tool Integration:** Investigated critical inference-latency bottlenecks in multi-turn, tool-augmented Large Language Models (LLMs); formulated an in-context tool composition mechanism within hierarchical skill abstractions, accelerating end-to-end response velocity by **2.0×** (from **20 s** to **10 s**).
- Topological Design of Agentic Execution Systems:** Conceptualized a novel hybrid architecture integrating heuristic routing, stateful workflows, and autonomous agentic fallbacks; conducted rigorous empirical analysis on its structural generalization and edge-case stability, yielding a state-of-the-art first-token latency of **5–6 s** under high-concurrency settings.

Ping An Technology — Semantics Laboratory (NLP Research Group)

12/2025 – 05/2026

Research Intern

- Mitigating Catastrophic Forgetting in Knowledge Alignment:** Developed an interactive conversational decision framework; fine-tuned parameterized LLMs (Qwen3-8B) via synthetic, chain-of-thought alignment topologies to systematically evaluate and mitigate *catastrophic forgetting* in resource-constrained models.
- Scalable High-Dimensional Representation Pipelines:** Engineered a robust dual-stage cascade infrastructure leveraging dense neural vector retrieval coupled with fine-grained LLM classification; expanded the evaluation space from **31** to **340** semantic classes while maintaining an empirical **Precision/Recall > 95%** across extreme long-tail class distributions.

SKILLS

Programming Python, C/C++, Java, CUDA C, Verilog, SQL, Lua

Machine Learning PyTorch, Hugging Face Ecosystem (Transformers, PEFT), Ray, DeepSpeed

LLM Infrastructure ReAct Framework, Tool-Use, MCP, Dify, SSE/AG-UI Streaming, Distributed Training

Data Systems Spark SQL, WeData, MySQL, Discrete-Event Simulation (SimPy)

Language TOEFL iBT **90** (R27 / W24 / L21 / S19); retake planned for August 2026 (targeting **105+**)